

Yet Another CHaracter Interaction sYstem Orchestrator: A Linearized Pipeline Architecture for Real-Time Embodied Conversational Agents

Yiyi Cai

Abstract

We present YACHIYO, a client-server architecture for real-time embodied conversational agents. The system models a complex multi-modal agent—comprising automatic speech recognition (ASR), large language model (LLM), retrieval-augmented generation (RAG), text-to-speech synthesis (TTS), and character animation—as a directed acyclic graph (DAG) of processing nodes, and linearizes it via topological sort into a sequential pipeline connected by thread-safe queues. We formalize *temporal consistency*—the guarantee that outputs respect causal, internal, and cancellation ordering—and show that linearization provides this guarantee with minimal implementation complexity. A dispatcher–receiver bracket recovers parallel execution for independent pipeline stages within the linear chain while preserving temporal consistency. A unified timestamp mechanism over WebSocket keeps the client and server synchronized, enabling streaming multi-modal output with low perceived latency. Each user owns an exclusive pipeline instance, while compute-heavy services run as shared standalone servers for independent scaling. We further introduce a WebRTC-based frame-level streaming mode that re-granularizes message boundaries from sentence-level to 20 ms frame-level, enabling continuous multi-modal output with sub-100ms synchronization. We describe the architecture, formalize the pipeline model, present latency measurements on a single NVIDIA RTX 5090 GPU, and discuss deployment experience.

1 Introduction

Problem. An embodied conversational agent must coordinate multiple processing stages in real time: speech recognition (ASR), language generation (LLM), information retrieval (RAG), speech synthesis (TTS), and character animation. Each stage has distinct computational requirements, latency characteristics, and software dependencies. The central challenge is: *how to orchestrate these heterogeneous stages into a system that is correct (outputs are ordered and synchronized), low-latency (the user perceives a responsive agent), scalable (multiple users can share expensive compute resources), and configurable (stages can be added, removed, or swapped without code changes)?*

Requirements. From first principles, a solution must satisfy:

1. **Temporal correctness:** outputs must arrive in the order implied by the dependency structure of the processing stages.
2. **Streaming:** downstream stages should begin processing as soon as partial results are available from upstream, rather than waiting for full completion.
3. **Client-server synchronization:** the rendering client and the processing server must agree on which output corresponds to which input, even under interruption and concurrency.
4. **Multi-modal synchronization:** audio, video, and metadata produced at different frame rates must be delivered to the client as a coherent, synchronized stream.

5. **Resource isolation with sharing:** each user needs an isolated execution context (conversation state, history), but expensive GPU models should be shared across users.
6. **Environment isolation:** different models may require incompatible software dependencies (e.g., different PyTorch versions for TTS vs. LLM).

Approach. We observe that the processing stages of a conversational agent form a directed acyclic graph (DAG) with a single entry and exit. We prove (Appendix A) that under mild assumptions, any such DAG can be *linearized*—reduced to a sequential chain of nodes connected by queues—without losing correctness. This linearization dramatically simplifies the execution model: each node is a thread consuming from one queue and producing to the next, with destination-based routing for skip-connections.

On top of this linearized pipeline, we layer:

- **Temporal consistency** (Section 4): we formalize three ordering properties (causal, internal, cancellation) and show that linearization guarantees all three with minimal implementation complexity.
- **Parallel branch execution** (Section 3.3): a dispatcher–receiver bracket recovers concurrent execution for independent nodes within the linear chain, using reverse-order emission and existing signal/forwarding mechanisms.
- **Intra-pipeline streaming:** the LLM streams sentences incrementally; each sentence immediately enters TTS, so the user hears audio while the LLM is still generating.
- **Timestamp-based synchronization:** a single timestamp per utterance propagates through the entire pipeline and back to the client, enabling correct ordering, stale-message filtering, and barge-in cancellation.
- **Frame-level re-granularization** (WebRTC mode): sentence-level messages are further split into 100 ms atomic groups of audio/video/data frames, synchronized via a GCD-based scheme.
- **Service-oriented compute:** heavy models run as standalone HTTP services with OpenAI-compatible APIs, called by lightweight per-user pipeline nodes.

We call the resulting system YACHIYO (**Y**et **A**nother **C**haracter Interaction **s**ystem **O**rchestrator). Existing frameworks address parts of this problem—LangChain [1] for LLM orchestration, LiveKit [4] for WebRTC transport—but none provides an integrated solution spanning pipeline linearization, temporal consistency, multi-modal streaming, client-server synchronization, and service-oriented scaling.

2 Related Work

Real-Time Conversational Agents. OpenAI’s GPT-4o Realtime API [2] provides end-to-end speech-to-speech generation but operates as a closed cloud service without user-controlled pipeline orchestration or custom model integration. OpenClaw [5], the most-starred open-source project on GitHub, is an autonomous AI agent platform that orchestrates multi-channel messaging, tool execution, and device integration through a WebSocket-based gateway. While it supports embodied interaction via paired device nodes, it focuses on general-purpose agent autonomy rather than real-time streaming pipeline orchestration for multi-modal conversational agents with character animation.

LLM Orchestration Frameworks. LangChain [1] provides composable abstractions for LLM chains including retrieval, tool use, and memory. LangGraph extends this with stateful graph-based workflows. However, these frameworks focus on text-level orchestration and do not address real-time streaming of audio, video, and animation data across client-server boundaries.

Node-Graph Processing Systems. ComfyUI [3] represents image and video generation workflows as visual DAGs where each node runs a model or transformation in-process. While flexible, this monolithic execution model requires all dependencies to coexist in one environment and does not support multi-user resource sharing or real-time streaming.

WebRTC for AI Agents. LiveKit Agents [4] and Daily Bots provide WebRTC infrastructure for connecting LLM-based agents to audio/video streams. These focus on transport-layer concerns and delegate pipeline logic to application code. Our system provides an integrated solution spanning pipeline orchestration, multi-modal synchronization, and real-time transport.

3 Pipeline Model

3.1 DAG Linearization

A pipeline is a JSON configuration listing processing nodes in topological order. Each node has a unique, monotonically increasing ID, a processing function (resolved from a module registry), and a configuration block specifying its data routing and downstream connections (see Appendix B for details and a complete example). Variable names are prefixed with the producing node’s ID (e.g., `3_text`) to avoid namespace collisions. Node IDs need not be contiguous, allowing insertion without renumbering.

The pipeline must guarantee *temporal consistency*: outputs corresponding to earlier inputs must arrive before outputs of later inputs, and outputs within a single utterance must arrive in production order. We formalize this property in Section 4; here we argue why linearization is the natural execution model.

Temporal consistency requires total ordering of messages. If messages can travel through the pipeline along multiple parallel paths, they may arrive at the output in a different order than they were produced. A merge point (where two branches rejoin) must resolve the ordering, which requires synchronization barriers—and barriers add complexity and latency.

A linear queue chain guarantees total order with minimal complexity. A chain of FIFO queues with single-threaded nodes preserves the order of messages by construction: if message m_1 enters before m_2 , then m_1 is dequeued and processed before m_2 at every node. No synchronization barriers or merge protocols are needed.

Alternative: full DAG parallel execution. Executing topologically parallel branches concurrently can improve throughput, but introduces merge points where messages from different branches must be reordered. Each merge point requires: (1) a synchronization barrier to wait for all incoming branches, (2) a merge protocol to establish the correct output order, and (3) handling of partial failures (one branch produces output while another is cancelled). For the same correctness guarantee, this is strictly more complex than the linear chain.

We prove in Appendix A that any DAG satisfying our assumptions can be linearized without loss of correctness.

Execution model. We restrict input DAGs to satisfy:

1. The graph has exactly one entry node and one exit node.
2. Node indices already form a valid topological ordering (edges always point from lower to higher IDs).
3. Each node’s output depends only on messages it has already received.

[DAG linearization diagram. Left: DAG with skip-connections. Right: linearized queue chain $Q_0 \rightarrow N_0 \rightarrow Q_1 \rightarrow N_1 \rightarrow \dots \rightarrow Q_n \rightarrow \text{send_queue}$, with destination-based routing shown]

Figure 1: DAG linearization. The topological sort produces a linear chain of nodes connected by queues. Messages carry a `destination` field; intermediate nodes forward messages not addressed to them.

At initialization, the system arranges nodes in topological order and connects them via a chain of thread-safe queues:

$$Q_{\text{in}} \rightarrow N_0 \rightarrow Q_1 \rightarrow N_1 \rightarrow \dots \rightarrow Q_n \rightarrow Q_{\text{send}}$$

Each node runs in its own thread, blocking on its input queue and pushing results to its output queue. Formally, each node N_i executes:

$$N_i : \quad m \leftarrow Q_i.\text{dequeue}(); \quad \begin{cases} Q_{i+1}.\text{enqueue}(f_i(m)) & \text{if } m.\text{dest} = i \\ Q_{i+1}.\text{enqueue}(m) & \text{otherwise (forward)} \end{cases} \quad (1)$$

where f_i is node i 's processing function, and $m.\text{dest}$ is the message's destination tag.

Trade-off: serialization overhead. If two nodes N_a, N_b are topologically parallel (no path from one to the other), linearization forces a sequential order. A message m destined for N_b must wait behind N_a 's processing of the preceding message. The added latency is bounded by:

$$\Delta L \leq \sum_{k \in \text{intervening}} L_k(m_{\text{prev}}) \quad (2)$$

where $L_k(m_{\text{prev}})$ is node k 's processing time for the preceding message. For independent nodes that both process the same message, reordering does not reduce the combined latency—the total is always $\sum L_k$ regardless of ordering. We show in Section 3.3 that this overhead can be reduced to $\max(L_k)$ via a dispatcher–receiver bracket without abandoning the linear chain.

3.2 Streaming within the Linear Pipeline

[Streaming timeline. Horizontal axis = time. Rows: LLM (streaming sentences $s_1, s_2, s_3, \dots$), TTS (synthesizing s_1 while LLM produces s_2), Client (playing s_1 audio while TTS synthesizes s_2). Show overlap between stages.]

Figure 2: Streaming pipeline timeline. The LLM streams sentences incrementally; each sentence is immediately synthesized by TTS and played by the client, overlapping in time.

Linearization does not imply that each stage must complete before the next begins. The LLM node streams tokens incrementally; a *StreamCutter* module segments them into sentences at punctuation boundaries. Each sentence is immediately pushed downstream as a separate

pipeline message carrying the *same* timestamp as the original utterance. TTS processes each sentence independently, and the client plays the resulting audio in FIFO order—guaranteed by temporal consistency (Section 4).

We formalize the latency. Let the LLM generate K sentences $s_1, s_2, \dots, s_K$. Let L_k^{LLM} be the time to generate sentence s_k (from the end of s_{k-1}), and L_k^{TTS} be the time to synthesize s_k .

In **batch mode**, the client receives audio only after all stages complete:

$$L_{\text{batch}} = L_{\text{ASR}} + \sum_{k=1}^K L_k^{\text{LLM}} + \sum_{k=1}^K L_k^{\text{TTS}} + L_{\text{net}} \quad (3)$$

In **streaming mode**, the client receives the first audio segment as soon as the first sentence is processed:

$$L_{\text{stream}}^{\text{first}} = L_{\text{ASR}} + L_1^{\text{LLM}} + L_1^{\text{TTS}} + L_{\text{net}} \quad (4)$$

Subsequent sentences overlap: while TTS synthesizes s_k , the LLM generates s_{k+1} . Defining $L_{K+1}^{\text{LLM}} = 0$ (no sentence after the last), the total time becomes:

$$L_{\text{stream}}^{\text{total}} = L_{\text{ASR}} + L_1^{\text{LLM}} + \sum_{k=1}^K \max(L_{k+1}^{\text{LLM}}, L_k^{\text{TTS}}) + L_{\text{net}} \quad (5)$$

When TTS dominates ($L_k^{\text{TTS}} \gg L_{k+1}^{\text{LLM}}$ for most k), the sum reduces to approximately $\sum_k L_k^{\text{TTS}}$, and total time approaches $L_{\text{ASR}} + L_1^{\text{LLM}} + \sum_k L_k^{\text{TTS}} + L_{\text{net}}$ —the LLM generation is almost entirely hidden behind TTS.

The speedup factor for first-output latency is:

$$S = \frac{L_{\text{batch}}}{L_{\text{stream}}^{\text{first}}} \quad (6)$$

3.3 Parallel Execution via Dispatcher–Receiver Bracket

Key observation. Linearization serializes independent nodes, adding latency ΔL (Section 3.1). However, this serialization is not inherent to the queue chain—it arises because each node must dequeue a message before the next can advance. If a message destined for a *later* node passes through an *earlier* node via forwarding (Equation 1, $\sim \mu\text{s}$) before the earlier node encounters its own message, both nodes process concurrently.

Mechanism: dispatcher and receiver. We introduce two meta-nodes that bracket a parallel region. Given k independent branch nodes $N_{b_1}, \dots, N_{b_k}$ (in topological order within the chain), the **dispatcher** and **receiver** operate as follows:

1. The dispatcher emits a **dispatch_start** signal (carrying upstream pass-through data) addressed to the receiver. This signal is forwarded through all branch nodes via destination-based routing.
2. The dispatcher emits k branch messages in *reverse topological order*: the message for N_{b_k} first, then $N_{b_{k-1}}$, down to N_{b_1} . Each message carries only the input fields required by its target branch.
3. The dispatcher emits a **dispatch_end** signal addressed to the receiver.

When N_{b_1} dequeues the branch messages, it first encounters the messages for $N_{b_2}, \dots, N_{b_k}$ (destinations do not match), forwarding each in μs . It then encounters its own message and begins processing. Meanwhile, the forwarded messages have already reached their respective nodes, which begin processing concurrently.

The receiver catches the **dispatch_start** and **dispatch_end** signals. Between them, it accumulates all arriving branch outputs. On **dispatch_end**, it merges the collected outputs with the pass-through data from **dispatch_start** and emits a single downstream message.

Why temporal consistency is preserved. FIFO ordering guarantees that all messages belonging to utterance u_i (with timestamp T_i) are dispatched, processed, and collected *before* any messages from u_{i+1} (with $T_{i+1} > T_i$). Specifically: (1) the dispatcher processes u_i before u_{i+1} (FIFO at the dispatcher’s input queue); (2) each branch node processes u_i ’s branch message before u_{i+1} ’s (FIFO at each branch’s queue); (3) the receiver collects u_i ’s `dispatch_end` before u_{i+1} ’s `dispatch_start` (FIFO at the receiver’s queue). Causal and internal order are therefore maintained. Cancellation consistency is preserved by the same timestamp-based filtering as in Section 4.4.

Latency improvement. For k parallel branches with processing times $L_1, \dots, L_k$, the serialized latency is $\sum_{i=1}^k L_i$, while the parallel latency is $\max(L_1, \dots, L_k)$. We evaluate this improvement on a motion generation pipeline in Section 7.

4 Temporal Consistency

The pipeline’s primary correctness requirement is *temporal consistency*: outputs delivered to the client must respect the ordering imposed by user utterances and the pipeline’s production sequence. We now formalize this property and show how linearization guarantees it.

4.1 Definition

Let $\{u_1, u_2, \dots\}$ be a sequence of user utterances with timestamps $T(u_1) < T(u_2) < \dots$. For each utterance u_i , the pipeline produces an ordered sequence of output messages $O_i = (o_{i,1}, o_{i,2}, \dots, o_{i,K_i})$. A pipeline is *temporally consistent* if it satisfies three properties:

1. **Causal order.** If $T(u_i) < T(u_j)$, then all outputs of u_i are delivered to the client before any output of u_j —unless u_i has been cancelled.
2. **Internal order.** Within a single utterance u_i , outputs are delivered in production order: $o_{i,1}$ arrives before $o_{i,2}$ before $\dots$ before o_{i,K_i} .
3. **Cancellation consistency.** After a cancellation event `cancel( $T$ )` is issued, no message with timestamp $< T$ is subsequently delivered to the client. The pipeline enforces this at every consuming node up to its exit; the client completes it end to end with its own flush — it initiated the cancel itself, or receives the server-originated one forwarded to it (Section 4.4).

4.2 Why Linearization Guarantees Temporal Consistency

Causal order. In a linearized pipeline, each node has exactly one input queue (FIFO) and processes messages one at a time in a single thread. If $T(u_i) < T(u_j)$, then u_i ’s messages enter the pipeline before u_j ’s. At every node in the chain, the FIFO queue guarantees that u_i ’s messages are dequeued and processed before u_j ’s. Therefore, u_i ’s outputs arrive at the send queue before u_j ’s.

Internal order. Within a single utterance, the LLM generates sentences $s_1, s_2, \dots, s_K$ sequentially, and the StreamCutter emits them in production order into the queue chain. Since every queue is FIFO and every node is single-threaded, s_1 ’s outputs are produced before s_2 ’s at every stage. The internal order is preserved end-to-end.

Linearization makes temporal consistency a structural guarantee rather than a runtime enforcement. The dispatcher–receiver bracket (Section 3.3) preserves this property by ensuring that all branch messages for u_i are dispatched, processed, and collected before any branch messages for u_{i+1} , as shown in the FIFO argument of Section 3.3.

4.3 Signals

Why signals are needed. A linearized pipeline processes *data* messages (audio, text, etc.) stage by stage. But some information is not data—it is *control*: “speech has started” (SoS), “speech has ended” (EoS), “recording has started/ended” (`recording_start/recording_end`). Control information often needs to reach the client faster than data, because the client must update its state (e.g., begin the listening animation) before any processed data arrives.

Why signals should not overtake data. If control signals could freely overtake data messages in the queue, ordering guarantees would break: the client might receive EoS before the last audio chunk. This is analogous to the hazard problem in CPU instruction pipelines, where out-of-order execution of control instructions (branches) relative to data instructions causes incorrect results.

Our design: declared routing, no drifting. Signals travel in the *same* queue as data, hop by hop, exactly like ordinary messages—there is no separate channel and no implicit “drifting”. What a node does with an arriving signal is declared *entirely in its pipeline config*, mirroring how `input_vars/pass_vars` declare field handling. Two declaration lists of explicit `{source, target}` entries—both fields always spelled out, same-name included; no shorthand and no implicit defaults—each list a strict one-to-one map (unique sources: a signal maps to exactly one name; unique targets: many-to-one renames would be an untraceable merge):

- `catch_signals`: deliver the signal to this node’s `process()` (renamed to its target, so a module’s internal handler names are decoupled from the wire names);
- `pass_signals`: relay the signal (renamed) to the next queue.

The four resulting states: *catch only* = consume and terminate; *catch + pass* = consume, then relay *after* `process()` returns, so any output the node emitted for the signal stays ahead of it (consume-then-relay); *pass only* = forward without processing, at queue-forwarding speed ($\sim\mu\text{s}$ per hop) rather than data-processing speed ($\sim\text{ms-s}$ per stage), without overtaking any queued message; *undeclared* = the signal dies at that node, dropped with a warning log. Every signal arriving at a node must therefore be declared, and the complete signal flow of a pipeline is readable from its config alone.

Signals use the same destination routing as data. A relayed (passed) copy is re-addressed to the node’s *first edge* (`next_nodes[0]`)—the same default destination a data output gets—so an ordinary signal such as SoS/EoS advances edge by edge, one declared `pass` hop at a time, checked at every node until the pipeline exit. A *directed* signal takes a non-default edge (e.g. the dispatcher’s `dispatch_start/end` to its receiver) and transits intermediate nodes via ordinary destination routing; at its target the same four states apply.

The emitting side is declared symmetrically: each module lists the signals it emits (by internal name) in its class contract `EMIT_SIGNALS`, and sends through `emit_signal()`; the wire name of every emission comes from the config’s `emit_signals` declaration — entirely config-defined with no contract default, the same rule as `catch_signals/pass_signals`, so a pipeline’s complete signal surface (emit, catch, pass) is readable from its config alone. Renaming emitted envelopes is what makes *nested* dispatcher/receiver brackets wireable: the inner pair emits, say, `inner_dispatch_start/end`, which the inner receiver catches, without colliding with the outer envelope. The dispatcher’s own signal semantics follow its logical mainline: its `pass_signals` relay *directed to the receiver* (transiting branch nodes, which neither see nor declare them), while `dispatch_signals`—the signal-side parallel of `dispatch_vars`—list which *caught* signals each branch receives as directed copies, by catch-target name (a signal must be caught to be dispatched; the same name may go to several branches). The validator enforces the exact match both ways: every catch target referenced by some branch list, every reference a real catch target.

Pipelines are validated statically at initialization time, and the validation is deliberately *per-node*: each module declares its own contract as class attributes (`REQUIRED_CATCH_SIGNALS` / `REQUIRED_INPUTS`, config-dependent variants via classmethods, plus `EMIT_SIGNALS` for what it sends), and the checking logic itself lives in the step classes (a pure classmethod `validate_config(config)` on the base class, extended via inheritance by structured modules such as the dispatcher), and the validator is pure dispatch over the nodes — declaration lists well-formed (explicit entries, one-to-one), catch targets equal to `required_catch_signals(config)` *exactly* (renamed wiring still satisfies the contract, since handlers match target names; an extra target would be silently swallowed by `process()`), `input_vars` targets equal to the module’s declared input set and `output_vars` sources equal to its product set — exactly, both directions (modules with config-defined lanes, such as the frame splitter’s data inputs or the frame collector’s demux outputs, mark that side free-form and only the fixed part is matched), `emit_signals` covering the module’s `EMIT_SIGNALS` exactly (both directions: every emission declared, every declaration a real emission), and node-internal reference consistency (dispatcher’s `dispatch_vars/dispatch_signals` must reference its own output/catch targets, the latter exactly both ways). On the wire side of any entry an explicit `null` is a declared opt-out rather than an omission: a null input source falls back to the module default, a null output target drops that product from the outgoing message, a null catch source leaves the slot unwired, and a null emit target suppresses the emission — the full contract stays visible in the config either way. Cross-node flow is intentionally *not* modeled: doing so would replicate each module’s routing semantics inside the validator and drift as modules evolve. Broken links between nodes instead surface at runtime through the four-state rules — an undeclared signal is dropped with a loud per-client log entry at the first node it reaches. Any static finding rejects the pipeline (HTTP 400 with details).

Environment failures are separated from configuration failures. Node initialization is sequential; each node’s `custom_init` runs under a timeout and records its outcome. A failed init (dependent model service down, bad settings entry) *fails fast*: nodes built so far are killed, the client state is reset to re-initializable, and the endpoint returns HTTP 503 with the failing node and cause — distinct from the validator’s 400 (fix the config) in that a 503 is retryable once the service is back.

The one exception requiring out-of-band delivery is cancellation, discussed in Section 4.4: cancellation is *not* a signal—its semantics are deliberately order-free (invalidate everything older than T), which is exactly why it may bypass the in-band ordering that signals must respect.

4.4 Timestamps and Cancellation

First-principles motivation. The causal origin of a conversational event is the client: the user spoke. The timestamp that identifies the event must therefore be assigned at a single point that observes client actions *in order* — before any stage that could reorder them. What corrupts ordering is not server-side assignment per se but assigning stamps at different points for related events; a single entry authority (the client itself over WebSocket, or the WebRTC gateway — whose DataChannel and media paths preserve arrival order — for a WebRTC session) keeps the stamps faithful to the causal order of user actions, and all comparisons stay within that one clock domain.

Design. Every message carries a `timestamp` field, minted outside the pipeline: the WebRTC gateway stamps everything it forwards with its own clock, and a direct WebSocket client supplies its own stamps. The entry validates identity on arrival — a numeric timestamp must be present (cancel included), and non-cancel timestamps must be monotonic (equal allowed, smaller rejected) — so every node downstream can trust the stamp without re-checking. Inside the pipeline, timestamps are only propagated and compared, never minted: a module stamps its outputs by copying identity from an upstream context (the `stamp` primitive), never from a clock. The timestamp is set once at utterance start and propagated unchanged through all pipeline

stages as a reserved field. The client uses it to associate responses with utterances, discard stale output, and maintain ordering under concurrency.

Why cancellation needs out-of-band delivery. When the user interrupts (barge-in), every pipeline node must *immediately* stop processing stale messages. If the cancel signal traveled through the data queue, it would wait behind the very messages it intends to invalidate—defeating its purpose. This is analogous to a CPU pipeline *flush*: upon a branch misprediction, all in-flight instructions must be discarded simultaneously, not sequentially.

Cancellation design. Each node has a dedicated `cancel_queue` (parallel to its data queue), checked before each dequeue operation. Every cancel carries a *source*: 0 for the pipeline boundary (client cancels and lifecycle teardown; the entry forces 0 on client submissions, so a client cannot impersonate a node), or the emitting node’s id (a pipeline node with event capability — e.g. a model-driven VAD detecting a speech onset — cancels the previous turn itself). A per-pipeline event handler dispatches each cancel to every node’s cancel queue *except the source* (the emitter has already handled its own state), and forwards server-originated cancels ($\text{source} \neq 0$) out to the client, whose own event handler mirrors the dispatch on its side. Each receiving node updates its local `cancel_timestamp`; a dequeued message *addressed to the node* with $\text{timestamp} < \text{cancel_timestamp}$ is silently dropped — the gate runs where a message is consumed, while forwarding hops relay unexamined.

The inequality is deliberately *strict*, uniformly across server and client. Issuers submit their semantic stamp (a barge-in cancel carries the very stamp of the turn it opens; teardown uses an infinite stamp to void everything); the event handler nudges every dispatched stamp down by a small epsilon — a purely numerical guard so the strict inequality reliably spares same-stamp messages across float round-trips, not a cancel offset.

The same queue carries a second control verb, `kill`: on pipeline teardown the server broadcasts `cancel( $\infty$ )` followed by `kill` to every node. The infinite stamp voids all content; `kill` makes the node exit its loop and run its cleanup hook. Long-running operations (streaming LLM and TTS calls) poll the control queue between chunks, so both verbs also abort work in flight, and a thread returning late from a blocking call still finds the verbs waiting in its own queue. The endpoint relays perform no CONTENT gating of their own: the inbound relay validates message identity only (timestamp presence and monotonicity, as above) and routes cancels to the control plane with source forced to 0; the outbound relay is a pure pass-through. Dropping decisions are entirely the nodes’ responsibility, so per-module drop policies remain possible.

Proof of no stale processing. After `cancel( $T$ )` is issued, consider any message m with $\text{timestamp}(m) < T$. At the moment the cancel is broadcast, m is in one of three states:

- **Waiting in a queue, not yet dequeued:** when the node dequeues m , it checks the cancel timestamp and drops m without processing. Every node’s watermark is current: non-source nodes receive the broadcast, and the excluded emitter aligns its own watermark to the cancel it minted before emitting it.
- **Currently being processed by node N_i :** the node completes (or aborts at a polling point) and any output $f_i(m)$ inherits m ’s timestamp; it is dropped when its *destination* node dequeues it. Intermediate forwarding hops relay it unexamined — the check runs where a message is consumed, so a stale message is dropped at its next consumer. If its destination is the pipeline exit, it leaves the pipeline and falls under the next case.
- **Already emitted by the final node:** the message has left the pipeline and rides transport hops that relay unexamined (the WebRTC gateway included); delivery-side handling belongs to the client, which flushes its own pipeline and playout — it initiated the cancel itself, or receives the forwarded server-originated one.

In the first two cases m (or any output derived from m) is never processed past the next node boundary after $\text{cancel}(T)$; the third case is outside the pipeline’s jurisdiction by construction, and its window is bounded by the exit-queue depth at the moment of cancellation. $\square$

5 WebRTC: Frame-Level Streaming

5.1 Two Paradigms: Event-Driven vs. Frame-Driven

WebRTC is a real-time communication protocol that provides *tracks* (continuous media streams for audio and video) and a *DataChannel* (a reliable or unreliable channel for arbitrary application data). The two communication modes in YACHIYO reflect fundamentally different timing paradigms:

WebSocket: event-driven, client-timestamped. Messages are produced only on events (user speech, LLM output, TTS completion); between events the channel is silent. The timing origin is the client’s wall clock, propagated through the pipeline.

WebRTC: frame-driven, server-timestamped. Audio and video tracks produce frames at fixed rates (50 fps, 30 fps) regardless of activity, emitting silence when idle. The timing origin is the server’s monotonic frame clock.

Paradigm adapters. Adapter nodes bridge the paradigms. The frame $\rightarrow$ event side is a pair: **FrameCollector** transforms each arriving frame group into per-lane pipeline messages (the group’s audio frames joined into one WAV chunk, video frames and data slots passed as lists), and **VAD** keeps a persistent ring of those audio chunks as rolling history; at **recording_start** it snapshots the pre-roll out of the ring into a per-segment buffer that the segment owns from then on, and finalizes the utterance at **recording_end** (with configurable post-roll) — a cancel then either voids the whole segment or leaves it untouched, never truncates it. The VAD can also drive itself from a detector service instead of client signals: a detected speech onset cancels the previous turn (a barge-in, dispatched through the event handler with the node as source) and opens the segment, a detected stop finalizes it, and client recording signals take manual control by suspending detection for that turn. The event $\rightarrow$ frame side is **FrameSplitter**, a clock-driven output module that splits TTS audio into 20 ms PCM frames paired with video and data. It overrides the standard message-driven processing loop with an absolute-time tick loop, pausing until a **connection_start** signal arrives and then outputting one synchronized group per tick regardless of whether speech data is available (filling with silence and idle frames when idle). The core pipeline remains purely event-driven.

5.2 Motivation: From Sentence-Level to Frame-Level

In WebSocket mode, the message boundary is a complete TTS segment ($\sim 0.5\text{--}1$ s of audio). The client receives a whole audio clip and then plays it, resulting in a perceptible “chunk” rhythm. WebRTC mode re-granularizes the output to frame-level:

- Audio: 20 ms PCM frames (960 samples at 48 kHz).
- Video: ~ 33 ms JPEG frames (30 fps).
- Data (text, action metadata): 50 ms slots (20 fps).

The output becomes a continuous stream rather than discrete chunks.

5.3 Multi-Modal Frame Synchronization

Three modalities with different frame rates must remain synchronized:

- Audio: $R_a = 48000/960 = 50$ fps

[Frame group diagram. One 100 ms group containing 5 audio frames (20 ms each), 3 video frames (33 ms each), and 2 data slots (50 ms each). $\text{GCD}(50, 30, 20) = 10 \rightarrow 100 \text{ ms group period.}$]

Figure 3: Multi-modal frame synchronization. Each 100 ms group bundles 5 audio frames, 3 video frames, and 2 data slots into an atomic synchronization unit.

- Video: $R_v = 30 \text{ fps}$
- Data: $R_d = 20 \text{ fps}$

Deriving the group period from first principles. We need the smallest period T_g such that each modality produces an *integer* number of frames within T_g . This requires $T_g \cdot R_a$, $T_g \cdot R_v$, and $T_g \cdot R_d$ to all be positive integers, i.e., $T_g = k / \text{gcd}(R_a, R_v, R_d)$ for the smallest positive integer $k = 1$. Thus:

$$R_g = \text{gcd}(R_a, R_v, R_d) = \text{gcd}(50, 30, 20) = 10 \text{ Hz}, \quad T_g = \frac{1}{R_g} = 100 \text{ ms} \quad (7)$$

This is the *finest* common synchronization granularity—any shorter period would require fractional frames in at least one modality. Each group contains $n_a = R_a/R_g = 5$ audio frames, $n_v = R_v/R_g = 3$ video frames, and $n_d = R_d/R_g = 2$ data slots.

The group is the atomic synchronization unit: the client consumes one group per T_g , dispatching exactly n_a audio frames, n_v video frames, and n_d data slots to their respective output channels. This guarantees that no modality drifts ahead of or behind others within a group. Per-sentence metadata (text, action, expression) rides a **meta** signal carrying the node’s pass-through fields; the group’s data slots are reserved for frame-aligned payloads (e.g. per-frame motion).

Video/data rates and video resolution are configurable per session and the GCD calculation adapts automatically (e.g., 24 fps video with 12 fps data yields $R_g = 2$, $T_g = 500 \text{ ms}$). The audio rate is pinned at 50 fps — WebRTC audio is 20 ms Opus frames on the wire and the SDP negotiates no framing — and video/data rates must divide the 90 kHz RTP clock within 10–60 fps; both constraints, along with the offered tracks matching the lanes the pipeline consumes and produces, are enforced when the connection is offered (HTTP 400 with the concrete mismatches).

5.4 Comparison: WebSocket vs. WebRTC Mode

Table 1: Comparison of communication modes.

Aspect	WebSocket	WebRTC
Message granularity	Sentence (~ 0.5 –1 s)	Frame group (100 ms)
Output continuity	Discrete chunks	Continuous stream
Audio format	Base64 WAV	Raw PCM (20 ms)
Synchronization	Timestamp per message	Atomic frame groups
Use case	Web UI, simple clients	2D avatar, immersive

6 System Design

6.1 Per-User Pipeline, Shared Services

First-principles motivation. Pipeline state consists of threads, queues, and per-user context (conversation history, session variables)—on the order of kilobytes per user. This is cheap to replicate, so each user gets a dedicated pipeline instance with full isolation. Model state, by contrast, consists of neural network weights loaded into GPU VRAM—on the order of gigabytes per model. Replicating model state per user is prohibitively expensive. The natural division is: replicate what is cheap (pipeline instances), share what is expensive (model services).

*[System architecture. Unity client ↔ WebSocket/WebRTC ↔ FastAPI server ↔
Standalone services (ASR, LLM, TTS, VectorDB)]*

Figure 4: System architecture. The client communicates with the pipeline server via WebSocket or WebRTC. The pipeline server orchestrates per-user pipeline instances, each calling shared standalone model services via HTTP APIs.

Each user gets an exclusive pipeline instance (threads, queues, conversation state). Compute-heavy models are deployed as standalone HTTP services:

- ASR: Qwen3-ASR server (GPU, vLLM backend)
- LLM: vLLM server (GPU)
- TTS: Qwen3-TTS server (GPU, faster-qwen3-tts)
- Vector DB: BGE embedding + FAISS server (GPU/CPU)

Pipeline nodes are thin HTTP clients: they create an OpenAI SDK client, call the API, and process the response. Multiple users’ pipelines share the same model servers concurrently. The model services expose OpenAI-compatible endpoints, so swapping a local model for a cloud API requires only a configuration change; services with no official API shape (the streaming VAD) use a small custom protocol behind the same settings mechanism.

6.2 Implementation Overview

The server is implemented in Python 3.12 using FastAPI and uvicorn for HTTP/WebSocket serving, aiortc for WebRTC, and the OpenAI Python SDK for model service calls. Each client follows the lifecycle: **register** → **init_pipeline** (from a JSON configuration) → **connect** (WebSocket or WebRTC) → pipeline runs → **unregister**. The pipeline is instantiated per-client with a dedicated thread per node and inter-node queues.

The client is built with Unity 6 (C#) using WebSocket-Sharp and Unity WebRTC for communication, and the Animator controller with blend shapes for character animation. The system comprises ~15 pipeline module types, dynamically discovered via a module registry. All configuration is JSON-based: pipeline configs, dataset configs, and model configs. Deployment uses per-service conda environments with a startup script that launches services sequentially and verifies port readiness.

7 Evaluation

All experiments are conducted on a single workstation equipped with an NVIDIA GeForce RTX 5090 GPU (32 GB VRAM). All local services—Qwen3-ASR-0.6B (ASR via vLLM), Qwen3.5-9B-AWQ (LLM via vLLM), Qwen3-TTS-0.6B via faster-qwen3-tts (TTS), BGE-M3 (vector

database)—run concurrently on this single GPU, occupying ~ 27 GB VRAM total. The test input is a 1.6-second Chinese speech clip.

7.1 Per-Stage and End-to-End Latency

Table 2: Per-stage and end-to-end latency: local pipeline vs. OpenAI API. All values in ms (mean $\pm$ std, 4 runs excluding warmup).

Stage	Local (RTX 5090)	OpenAI API
ASR	28 ± 1	886 ± 386
LLM (total generation)	290 ± 88	1252 ± 189
TTS (first sentence)	875 ± 323	1788 ± 454
Server first audio	1101 ± 331	3505 ± 860
E2E first audio (client)	1139 ± 347	3541 ± 855
E2E total	2228 ± 754	7264 ± 2927

The local pipeline achieves a first-audio latency of ~ 1.1 s, with TTS synthesis dominating (~ 0.9 s per sentence). The OpenAI API pipeline incurs ~ 3.5 s first-audio latency, primarily due to network round-trip times. The streaming architecture means that the total response time (~ 2.2 s local) includes overlapping computation: while TTS synthesizes sentence n , the LLM is already generating sentence $n+1$.

7.2 Multi-User Scalability

Table 3: Multi-user latency on a single RTX 5090 (local pipeline). All values in milliseconds.

Users	First Audio (avg)	First Audio (max)	Total (avg)
1	966	966	1372
2	2070	2072	3587
3	2878	2883	5427

Table 3 reports a separate multi-user experiment (single run per configuration); the 1-user baseline differs slightly from Table 2 (which averages 4 runs) due to natural variance.

When U concurrent users share services on a single GPU, latency depends on the serialization behavior of each service. The LLM server (vLLM) supports continuous batching and can partially parallelize concurrent requests, but the TTS server (Qwen3-TTS) processes requests sequentially. Modeling the TTS as the serialization bottleneck, the expected first-audio latency for user u (ordered by arrival) is approximately:

$$\mathbb{E}[L_u] \approx L_{\text{stream}}^{\text{first}} + (u - 1) \cdot L_{\text{bottleneck}} \quad (8)$$

where $L_{\text{bottleneck}}$ is the processing time of the slowest shared service (TTS in our case, ~ 0.8 s). The average across all users is:

$$\bar{L} \approx L_{\text{stream}}^{\text{first}} + \frac{U - 1}{2} \cdot L_{\text{bottleneck}} \quad (9)$$

Table 3 confirms this model: for 1–3 users, first-audio latency increases by ~ 950 ms per additional user, consistent with TTS serialization (~ 875 ms plus scheduling overhead). With R replicas of the bottleneck service, the effective capacity scales as $\bar{L} \approx L_{\text{stream}}^{\text{first}} + \frac{U-1}{2R} \cdot L_{\text{bottleneck}}$.

7.3 Configurability

The same pipeline framework supports multiple configurations without code changes:

Table 4: Example pipeline configurations.

Config	Pipeline
Conversation	ASR $\rightarrow$ LLM $\rightarrow$ RAG $\times 2 \rightarrow$ TTS
Conversation (WebRTC)	AudioCollector $\rightarrow$ ASR $\rightarrow$ LLM $\rightarrow$ RAG $\times 2 \rightarrow$ TTS $\rightarrow$ FrameSplitter
Conversation (cloud)	ASR ^{cloud} $\rightarrow$ LLM ^{cloud} $\rightarrow$ RAG $\times 2 \rightarrow$ TTS ^{cloud}
Motion gen. (parallel)	ASR $\rightarrow$ LLM $\rightarrow$ RAG $\rightarrow$ Dispatch $\rightarrow$ MotionGen $\parallel$ TTS $\rightarrow$ Receive
VTuber livestream	DanmakuBuffer $\rightarrow$ LLM $\rightarrow$ RAG $\rightarrow$ Dispatch $\rightarrow$ MotionGen $\parallel$ TTS $\rightarrow$ Receive

Switching between local and cloud models requires only editing the model JSON configuration file (changing `api_base` and `api_key`).

7.4 Motion Generation and Parallel Execution

An alternative pipeline replaces RAG with a text-to-motion generation service (HY-Motion), which produces SMPLH body parameters from LLM action descriptions. The HY-Motion per-stage latency (mean $\pm$ std over 3 runs): prompt rewriting 127 ± 10 ms, text encoding 53 ± 1 ms, ODE solver (50 steps) 276 ± 5 ms, decode + SMPLH 9 ± 0 ms; end-to-end 884 ± 36 ms.

Since MotionGen and TTS are independent (both consume LLM output), they can be parallelized via the dispatcher–receiver bracket (Section 3.3):

Table 5: First-audio latency: sequential vs. parallel motion generation pipeline (mean $\pm$ std over 3 runs, single RTX 5090).

Configuration	First Audio (ms)
Sequential (MotionGen $\rightarrow$ TTS)	1649 ± 199
Parallel (MotionGen $\parallel$ TTS)	1536 ± 172
Improvement	113 ms ($\downarrow 6.9\%$)

The improvement (113 ms, 6.9%) is modest because the MotionGen service adds ~ 540 ms to the pipeline (sequential first audio 1649 ms vs. standard 1110 ms), but TTS (~ 875 ms) still dominates the critical path in the parallel configuration. The theoretical maximum improvement equals the MotionGen latency (~ 540 ms), but is limited by the dispatcher–receiver overhead and the fact that both branches share GPU resources.

8 Discussion and Conclusion

Limitations. The dispatcher–receiver bracket (Section 3.3) supports fork–join parallelism at any point in the pipeline, but requires that parallel branches be independent (no data flow between branches within a bracket). Nested brackets (parallel regions within parallel regions) are possible in principle but have not been tested. The WebRTC FrameSplitter assumes fixed frame rates; adaptive rate adjustment would be needed for variable-bandwidth networks.

Future Work. Planned extensions include: (1) nested dispatcher–receiver brackets for multi-level parallelism; (2) vision input modality (camera $\rightarrow$ image understanding); (3) adaptive frame rate synchronization for WebRTC; (4) multi-GPU TTS serving for concurrent multi-character support.

Conclusion. We presented YACHIYO, a linearized pipeline architecture for real-time embodied conversational agents. The key contributions are:

1. A DAG-to-linear-chain transformation that ensures correct execution order while enabling streaming at every stage, reducing first-output latency from the sum of all stage latencies to the critical path of a single sentence.
2. A formal definition of *temporal consistency* (causal order, internal order, cancellation consistency) and a proof that the linearized pipeline guarantees all three properties.
3. A dispatcher–receiver bracket that recovers parallel execution within the linear chain by exploiting reverse-order emission and destination-based forwarding, while preserving temporal consistency.
4. Two levels of streaming granularity: sentence-level (WebSocket) and frame-level (WebRTC), with a GCD-based multi-modal synchronization mechanism for atomic 100 ms frame groups.
5. Timestamp-based client-server synchronization enabling correct ordering, stale message filtering, and clean barge-in cancellation.
6. Separation of per-user lightweight pipeline instances from shared GPU-hosted model services, connected via an OpenAI-compatible API layer that allows transparent swapping between local and cloud backends.

On a single NVIDIA RTX 5090 GPU hosting all services, the system achieves ~ 1.0 s first-audio latency for a fully local pipeline and supports up to 3 concurrent users within interactive latency bounds (< 3 s first audio).

References

- [1] Chase, H. LangChain: Building applications with LLMs through composability. <https://github.com/langchain-ai/langchain>, 2022.
- [2] OpenAI. GPT-4o and the Realtime API. <https://platform.openai.com/docs/guides/realtime>, 2024.
- [3] comfyanonymous. ComfyUI: A powerful and modular stable diffusion GUI and backend. <https://github.com/comfyanonymous/ComfyUI>, 2023.
- [4] LiveKit Inc. LiveKit Agents: Build real-time AI agents. <https://docs.livekit.io/agents/>, 2024.
- [5] OpenClaw. OpenClaw: Personal AI assistant you run on your own devices. <https://github.com/openclaw/openclaw>, 2025.

Appendix

A Proof: DAG Linearization

We prove that the class of DAGs arising from conversational agent pipelines can be linearized into a sequential chain without loss of correctness.

Definitions. Let $G = (V, E)$ be a directed acyclic graph where $V = \{v_1, v_2, \dots, v_n\}$ are processing nodes and $E \subseteq V \times V$ are data dependencies. We assume:

- (A1) G has a unique source v_1 (no incoming edges) and a unique sink v_n (no outgoing edges).
- (A2) Every edge $(v_i, v_j) \in E$ satisfies $i < j$ (i.e., node indices already form a valid topological ordering). This is without loss of generality—nodes can always be renumbered.
- (A3) Each node’s output depends only on messages it has already received (causal processing).

Theorem. Under assumptions (A1)–(A3), G can be executed correctly by a linear chain $v_{\pi(1)} \rightarrow v_{\pi(2)} \rightarrow \dots \rightarrow v_{\pi(n)}$ connected by FIFO queues, where π is any topological ordering of G .

Proof.

1. **Existence of valid ordering.** By (A2), node indices already form a topological ordering: for every edge $(v_i, v_j) \in E$, we have $i < j$. We use this index ordering directly as the linearization order $\pi = \text{id}$.
2. **Message routing.** Arrange nodes in topological order and connect consecutive nodes by a single FIFO queue:

$$Q_0 \rightarrow v_{\pi(1)} \rightarrow Q_1 \rightarrow v_{\pi(2)} \rightarrow Q_2 \rightarrow \dots \rightarrow Q_{n-1} \rightarrow v_{\pi(n)} \rightarrow Q_n$$

When v_i produces a message destined for v_j (where $j > i$ in topological order), the message is tagged with **destination** = j and placed in Q_i .

3. **Forwarding.** Each node v_k dequeues a message from Q_{k-1} . If **destination** $\neq k$, v_k forwards the message unchanged to Q_k without processing. If **destination** = k , v_k processes the message.
4. **Correctness.** In the linearized pipeline, each node has exactly one input queue and processes messages one at a time. By the topological ordering guarantee, v_j will eventually dequeue the message because:
 - The message enters the chain at position $i < j$.
 - Every intermediate node v_k ($i < k < j$) forwards it because **destination** = $j \neq k$.
 - FIFO ordering within each queue preserves the temporal order of messages from the same source.

By (A3), v_j can correctly process the message because all its dependencies (from nodes earlier in the topological order) have already been processed and their outputs forwarded. For multi-dependency nodes, all predecessors are placed earlier in the linearization, so their outputs arrive at v_j in the correct order. By (A3), v_j ’s output depends only on messages it has already received, so it can accumulate inputs from predecessors via internal state and produce its output upon receiving the final dependency—no external synchronization is required.

Trade-off: serialization of parallel nodes. If v_a and v_b are topologically parallel (no path from one to the other), linearization forces one to precede the other. A message destined for v_b must wait in the queue behind v_a ’s processing. The added latency is:

$$\Delta L = \sum_{k \in \text{intervening nodes}} L_k$$

where L_k is the processing time of node k for the preceding message.

For independent nodes that both process the same message, reordering does not reduce the combined latency—the total is always $\sum L_k$. The dispatcher–receiver bracket (Section 3.3) reduces this to $\max(L_k)$.

B Data Routing and Configuration

Each inter-node message is a flat key-value object. Variable names are prefixed with the producing node’s ID to avoid collisions (e.g., `3_text`). Three routing mechanisms handle data flow without requiring nodes to know the full pipeline topology:

- **input_vars**: maps incoming fields to the internal names the module reads (explicit `{source, target}` entries, one-to-one).
- **output_vars**: maps the module’s products to outgoing field names.
- **pass_vars**: carries metadata unchanged across nodes that do not consume it (e.g., the LLM’s action tag passes through TTS to reach the client).
- **catch_signals** / **pass_signals** / **emit_signals**: the signal counterparts (Section 4.3); every signal arriving at a node must be declared as caught and/or passed, so the listing’s downstream nodes declare **pass_signals** for the LLM’s SoS/EoS.

Var declarations follow the same exact per-node contract as signals: **input_vars** targets must equal the module’s declared input set and **output_vars** sources its product set, both directions, and an explicit `null` on the wire side is a declared opt-out — a null input source uses the module default, a null output target keeps that product off the wire (visible in the listing’s **language**, **raw_text**, **speaker** and **duration** entries).

Pass-through is only applied when a node actually processes a message; forwarded messages (destination mismatch) are passed unchanged. A **destination** field enables skip-connections to non-adjacent nodes; the special id `-1` denotes the pipeline exit, so terminal nodes declare **next_nodes**: `[-1]`. Node IDs need not be contiguous but must be monotonically increasing.

Listing 1: Pipeline configuration (valid as written).

```

1 {
2   "pipeline": [
3     {
4       "node_id": 1,
5       "function": "call_openai_asr",
6       "config": {
7         "input_vars": [{"source": "audio_file",
8                           "target": "audio_file"}],
9         "output_vars": [{"source": "result",
10                           "target": "1_text"},
11                           {"source": "language",
12                             "target": null}],
13         "next_nodes": [3],
14         "model": "qwen_asr"
15       }
16     },
17     {
18       "node_id": 3,
```

```

19     "function": "call_openai_llm",
20     "config": {
21         "input_vars": [{"source": "1_text",
22                         "target": "prompt"}],
23         "output_vars": [{"source": "text",
24                           "target": "3_text"},
25                          {"source": "raw_text",
26                           "target": null}],
27         "emit_signals": [{"source": "SoS",
28                             "target": "SoS"},
29                           {"source": "EoS",
30                            "target": "EoS"}],
31         "next_nodes": [6],
32         "model": "qwen"
33     }
34 },
35 {
36     "node_id": 6,
37     "function": "call_openai_tts",
38     "config": {
39         "input_vars": [{"source": "3_text",
40                         "target": "text"},
41                          {"source": null,
42                           "target": "language"},
43                          {"source": null,
44                           "target": "speaker"},
45                          {"source": null,
46                           "target": "duration"}],
47         "output_vars": [{"source": "audio_file",
48                           "target": "audio_data"},
49                          {"source": "duration",
50                           "target": null}],
51         "pass_signals": [{"source": "SoS",
52                             "target": "SoS"},
53                           {"source": "EoS",
54                            "target": "EoS"}],
55         "next_nodes": [-1],
56         "model": "qwen_tts"
57     }
58 }
59 ]
60 }

```

Note: `node_id` values (1, 3, 6) are non-contiguous but monotonically increasing, allowing future insertion of new nodes (e.g., node 2 or 4) without renumbering existing nodes.

C Implementation Details

C.1 OpenAI-Compatible API Layer

Motivation. Pipeline nodes communicate with model services via HTTP, which provides environment isolation (each service runs in its own process with its own dependencies). The remaining design choice is the HTTP protocol: a custom API or an existing standard. The OpenAI API has become the de facto standard in the LLM ecosystem—every major model serving framework (vLLM, Ollama, TGI) already implements it, and client SDKs are freely available. Adopting this standard minimizes integration effort: swapping a local model for a cloud API (or vice versa) requires only changing the configuration file, with no code modification.

All standalone services expose OpenAI-compatible endpoints:

- ASR: `POST /v1/audio/transcriptions` (Whisper API format)
- LLM: `POST /v1/chat/completions` (Chat Completions API format)
- TTS: `POST /v1/audio/speech` (TTS API format)

Model selection is config-driven: each model has a JSON configuration file specifying `api_base`, `api_key`, `model`, and `extra` parameters. Each module makes an initialization call at startup to verify service availability and trigger any needed model loading (e.g., TTS voice model switching).

C.2 Comparison with Monolithic Execution

Systems like ComfyUI [3] execute all nodes within a single runtime, loading every model into the same process and GPU memory space. While this avoids inter-process communication overhead, it creates several constraints that conflict with the requirements of a multi-user, real-time conversational system:

1. **Resource sharing:** in a monolithic system, each session must load its own model instances or implement complex in-process model multiplexing. In our architecture, a single GPU-hosted model server serves many concurrent users transparently, with the serving framework (e.g., vLLM’s continuous batching) handling request scheduling.
2. **Independent scaling:** if the TTS service becomes the bottleneck, it can be replicated to additional GPUs without touching the pipeline or other services. In a monolithic system, scaling one component requires replicating the entire runtime.
3. **Environment isolation:** each service runs in its own environment (e.g., different PyTorch versions for TTS and LLM), and swapping a model from local to cloud is a configuration edit.
4. **Fault isolation:** a crash in one service does not bring down the entire system; other users’ pipelines continue operating.

The trade-off is HTTP overhead per call ($\sim 1\text{--}10\text{ ms}$), which is negligible when per-node inference time is $30\text{ ms--}1\text{ s}$.

C.3 Client Architecture

The Unity client mirrors the server’s pipeline design: modules (recording, content, action, audio playback, WebSocket/WebRTC communication) are connected in a sequential queue chain with the same destination-based routing (per-module edge tables sharing the -1 exit-vertex convention), timestamp propagation, and cancel-queue mechanism. Signals follow the same four-state declared routing as Section 4.3: each module’s catch/pass lists are serialized scene fields—the client’s analogue of the server’s pipeline config—with the same consume-then-relay ordering, and undeclared signals dropped loudly. Renaming is deliberately omitted on the client: its chain is short and single-instance, so wire names are simply the module contract names. Contracts mirror the server’s per-node validation philosophy: each module declares the signals it cannot function without (`RequiredCatchSignals`), checked at initialization strictly against the module’s own wiring, and any finding or module init failure stops the pipeline with a `pipeline_error` event—the client-side counterpart of the server’s 400/503 split. Two boundary conventions complete the picture: the network modules’ catch lists serve as the per-scene *export manifest* (declaring which client-side signals cross into the remote server pipeline), and after a successful start the client injects a `pipeline_ready` message addressed to the exit vertex, whose traversal of every queue doubles as a liveness self-test. This symmetry means the client and server share the same correctness guarantees (temporal consistency, cancellation consistency) without additional coordination logic.

A three-state machine governs the interaction flow:

- **Idle:** character plays idle animation, awaiting a `listening_start` activation.

[State machine. Three states: Idle, Listening, Answering. Transitions: Idle $\xrightarrow{\text{listening_start}}$ Listening $\xrightarrow{\text{listening_end}}$ Answering $\xrightarrow{\text{answering_end}}$ Idle. Barge-in: Answering $\xrightarrow{\text{listening_start}}$ Listening. Dashed arrow: any state $\xrightarrow{\text{cancel}}$ Idle.]

Figure 5: Client-side state machine governing agent behavior.

- **Listening:** dual-threshold VAD detects speech; audio is recorded and sent to the server. On `listening_end` it advances to Answering.
- **Answering:** the client plays back server responses (audio, text, animation) in FIFO order. On `answering_end` it returns to Idle; a barge-in (`listening_start`) interrupts playback and returns to Listening.

The state machine is purely client-side; the server does not track client state. For motion generation pipelines, retargeting happens on the server: the motion node converts SMPL-H results to an engine-native humanoid format (per-bone rotations keyed by Unity’s `HumanBodyBones`, plus root motion). The client’s humanoid motion player only consumes this pre-retargeted format, playing back a continuous frame buffer at the motion’s native framerate with crossfade blending between motions and idle auto-refill.

C.4 LLM Workflow

The LLM node internally implements a structured dialogue management workflow inspired by the SillyTavern character card format:

Multi-turn history is persisted per client. Character personality is defined in external “lore-book” files whose entries are activated by keyword or vector similarity and injected into the prompt at each turn. The LLM streams tokens, segmented into sentences by the StreamCutter; action tags are parsed for downstream animation mapping; tool calls are executed in a loop. This workflow runs entirely within the LLM node—upstream and downstream nodes are unaware of it.

C.5 Backpressure

Each inter-node queue accepts an optional `max_queue_size` parameter. When a producer attempts to enqueue into a full queue, the `put()` call blocks until the consumer drains a slot. The send queue supports a separate `send_queue_max_size` for the same purpose.

C.6 BaseProcessingStep

Every pipeline node extends `BaseProcessingStep`, which implements the core execution loop described in Section 3. Each instance runs in a dedicated thread and provides:

1. **Queue-driven run loop:** blocks on the input queue, processes one message at a time, and enqueues results to the output queue.
2. **Destination routing:** checks each message’s `destination` field; forwards non-matching messages to the output queue without processing (Section 3.1).
3. **Signal handling:** each node declares `catch_signals` / `pass_signals` in its pipeline config (renaming supported, like the var declarations); caught signals go to `process()`, passed signals are relayed at forwarding speed, undeclared signals are dropped with a warning (Section 4.3).

4. **Control plane:** drains the control queue before each dequeue; messages addressed to the node with `timestamp < cancel_timestamp` are silently dropped, and the `kill` verb exits the loop through the node's cleanup hook (Section 4.4).
5. **Variable routing:** extracts input fields via `input_vars`, attaches output fields via `output_vars`, and carries metadata via `pass_vars` (Appendix B).

Subclasses override a single `process()` method containing the node-specific logic. The base class handles all pipeline mechanics, so module authors need not reason about routing, ordering, or cancellation.

C.7 SpanProcessingStep

Some modules process data spanning multiple messages (e.g., the VAD's segment accumulates audio from `recording_start` until `recording_end`). `SpanProcessingStep` extends `BaseProcessingStep` with: `start_span(timestamp)` to begin accumulation, `end_span()` to complete it, and `on_span_cancel()` to discard accumulated state on cancellation.

C.8 Connection Lifecycle

When a WebSocket or WebRTC connection closes, the server automatically disposes the client's pipeline: it broadcasts `cancel( $\infty$ )` followed by `kill` into every node's control queue, so in-flight work aborts at its polling points and each thread exits through its cleanup hook. For WebRTC, a `connection_start` signal is sent upon WebSocket establishment, allowing clock-driven modules to begin output only after the transport is ready.